

SIDEKIT – Python Toolkit for Speaker Recognition

Presented by:

Kong Aik LEE

Institute for Infocomm Research (I²R)
A*STAR, Singapore

Overview

- Recap and overview
- Training a TVM using SIDEKIT
 - FeatureExtractor, FeatureServer, StatServer
 - Whitening
 - Minimum divergence (MD) re-estimation
- Probabilistic LDA (PLDA)
 - Tying across frame
 - Simplified PLDA
 - Speaker adaptation
- Interesting research topics

SIDEKIT – Python Toolkit for Speaker Recognition

RECAP AND OVERVIEW

I-vector extraction (1/3)

- **Compression** process – an i-vector is a fixed-length low-dimensional representation of a variable-length speech utterance [Dehak et al, 2011].
- **MAP estimate** – posterior mean of the latent variable \mathbf{h} in a multi-Gaussian factor analysis model (i.e., total variability model)

i-vector \nearrow

$$\phi = \underset{\mathbf{h}}{\operatorname{argmax}} \left[\prod_{c=1}^C \prod_{t=1}^{N_c} \mathcal{N}(o_t | \boldsymbol{\mu}_c + \mathbf{W}_c \mathbf{h}, \boldsymbol{\Phi}_c) \right] p(\mathbf{h})$$

$$\phi = \underbrace{\left(\sum_{c=1}^C N_c \mathbf{W}_c^T \boldsymbol{\Phi}_c^{-1} \mathbf{W}_c + I \right)}_{\mathbf{L}^{-1}} \sum_{c=1}^C \mathbf{W}_c^T \boldsymbol{\Phi}_c^{-1} \underbrace{\left(\sum_{t=1}^{N_c} \gamma_t (o_t - \boldsymbol{\mu}_c) \right)}_{\text{First-order statistics (centred)}}$$

Zero-order statistics \nearrow


$$N_c = \sum_{t=1}^T \underbrace{\gamma_t(c)}_{\text{Alignment of frame } o_t \text{ to Gaussian } c}$$

$$\gamma_t(c) = \frac{\omega_c \mathcal{N}(o_t | \boldsymbol{\mu}_c, \boldsymbol{\Phi}_c)}{\sum_{j=1}^C \omega_j \mathcal{N}(o_t | \boldsymbol{\mu}_j, \boldsymbol{\Phi}_j)}$$

Alignment of frame o_t to Gaussian c

I-vector extraction (2/3)

- Frame alignment with *deep neural network* (DNN)
 - NN trained for phone state classification was used to align the frames.
 - Each output node of the DNN is trained to estimate the posterior probability of tied-states given the acoustic observations.



I-vector extraction (3/3)

- Vector notation

$$\phi_r = \left(\sum_{c=1}^C N_{r,c} \mathbf{W}_c^T \mathbf{\Phi}_c^{-1} \mathbf{W}_c + \mathbf{I} \right)^{-1} \left(\sum_{c=1}^C \mathbf{W}_c^T \mathbf{\Phi}_c^{-1} \sum_{t=1}^{N_{r,c}} (o_t - \boldsymbol{\mu}_c) \right)$$
$$\phi_r = \underbrace{(\mathbf{I} + \mathbf{T}^T \mathbf{\Sigma}^{-1} \mathbf{N}_r \mathbf{T})^{-1}}_{\mathbf{L}_r^{-1}} \cdot \mathbf{T}^T \mathbf{\Sigma}^{-1} (\mathbf{F}_r - \mathbf{N}_r \mathbf{m}_o)$$

- \mathbf{F}_r is a $CD \times 1$ vector obtained by stacking the first-order statistics from all components. \mathbf{N}_r is a $CD \times CD$ diagonal matrix consisting of diagonal blocks $N_{r,c} \times \mathbf{I}$

$$\mathbf{F}_r = \begin{bmatrix} \mathbf{F}_r(1) \\ \mathbf{F}_r(2) \\ \vdots \\ \mathbf{F}_r(C) \end{bmatrix} \quad \mathbf{N}_r = \begin{bmatrix} N_{r,1} \times \mathbf{I} & 0 & \cdots & 0 \\ 0 & N_{r,2} \times \mathbf{I} & 0 & \vdots \\ \vdots & 0 & \ddots & 0 \\ 0 & \cdots & 0 & N_{r,C} \times \mathbf{I} \end{bmatrix}$$

Total variability model

- The **total variability matrix** \mathbf{T} is obtained by stacking up the loading matrices from each Gaussian:

Total variability matrix

$$\mathbf{T} = \begin{bmatrix} \mathbf{W}_1 \\ \mathbf{W}_2 \\ \vdots \\ \mathbf{W}_C \end{bmatrix}$$

UBM mean vectors and covariance matrices

$$\mathbf{m}_o = \begin{bmatrix} \boldsymbol{\mu}_1 \\ \boldsymbol{\mu}_2 \\ \vdots \\ \boldsymbol{\mu}_C \end{bmatrix}$$

$$\boldsymbol{\Sigma} = \begin{bmatrix} \boldsymbol{\Phi}_1 & 0 & \cdots & 0 \\ 0 & \boldsymbol{\Phi}_2 & 0 & \vdots \\ \vdots & 0 & \ddots & 0 \\ 0 & \cdots & 0 & \boldsymbol{\Phi}_C \end{bmatrix}$$

- Total variability model

$$p(o_t | c) = \mathcal{N}(o_t | \boldsymbol{\mu}_c + \mathbf{W}_c \mathbf{h}_r, \boldsymbol{\Phi}_c) \quad \text{for } c = 1, 2, \dots, C$$

$$\boldsymbol{\mu}_{r,c} = \boldsymbol{\mu}_c + \mathbf{W}_c \mathbf{h}_r \quad \text{for } c = 1, 2, \dots, C$$

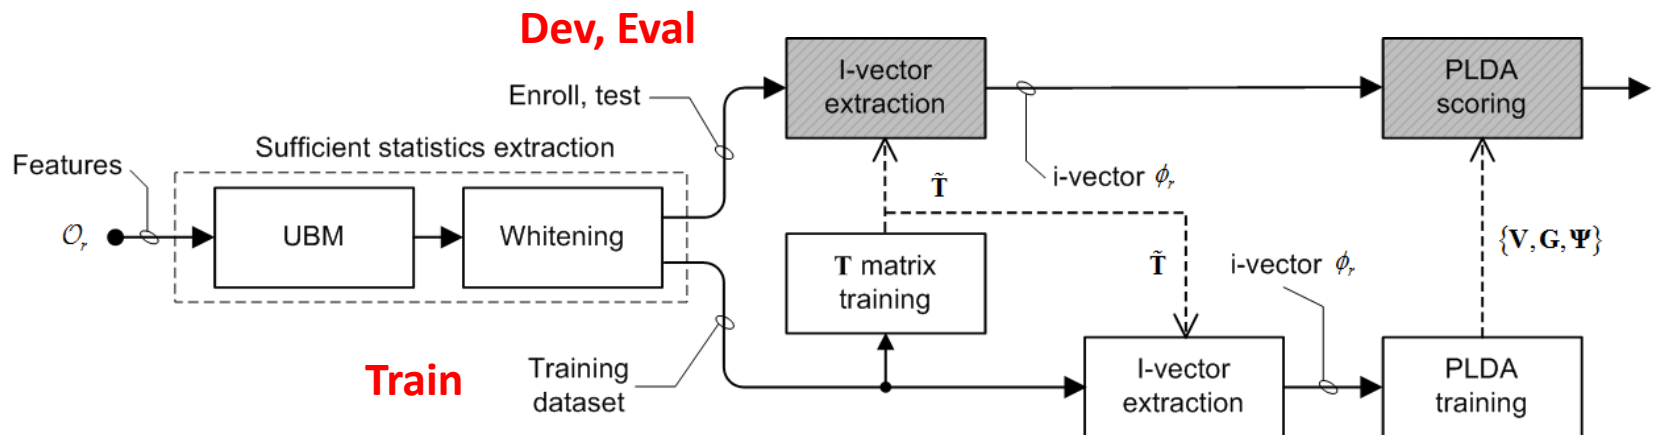
mean supervector

$$\mathbf{m}_r = \mathbf{m}_o + \mathbf{T} \mathbf{h}_r$$

Session index, r

I-vector PLDA pipeline

- I-vector extraction followed by PLDA scoring



[K. A. Lee and H. Li, Interspeech 2017]

- PLDA scoring

$$s(\phi_t, \phi_e) = \frac{p(\phi_t, \phi_e)}{p(\phi_t)p(\phi_e)} = \frac{p(\phi_t | \phi_e)p(\phi_e)}{p(\phi_t)p(\phi_e)} = \frac{p(\phi_t | \phi_e)}{p(\phi_t)}$$

Test
Enroll

SIDEKIT – Python Toolkit for Speaker Recognition

TRAINING A TVM USING SIDEKIT

SIDEKIT

- End-to-end python implementation from feature extraction, scoring, calibration (Python BOSARIS)
- Main classes: FeatureExtractor, FeatureServer, StatServer
- Reference:

A. Larcher, K. A. Lee, and S. Meignier, “An extensible speaker identification SIDEKIT in Python,” in *Proc. IEEE ICASSP*, 2016, pp. 5095 – 5099.

- Tutorial page:

<http://www-lium.univ-lemans.fr/sidekit/overview/index.html>

- PIP (with Anaconda and virtual environment)

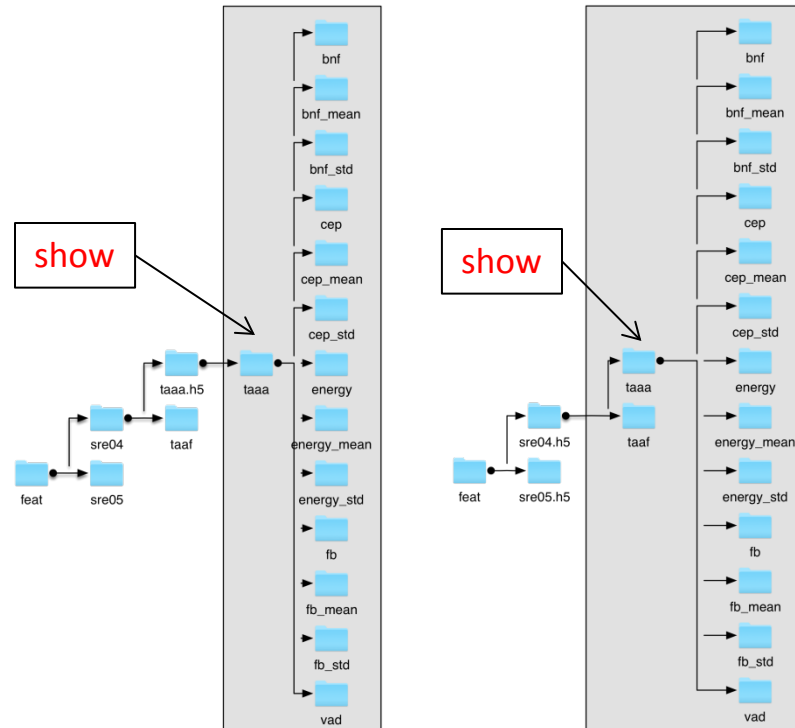
```
$ pip install sidekit
```

Feature Extractor (1/2)

- `FeatureExtractor` – a Python class providing functionalities to
 - Extract features (MFCC, LFCC) from audio files (sph, wav, raw)
 - Energy-based voice activity detection (VAD)
 - Features (e.g., MFCC, VAD, log energy etc.) are stored as datasets in hdf5 format

- Hierarchical architecture of hdf5 allows

- One hdf5 feature file per audio recording
- One hdf5 feature file for a collection of audio recordings



Feature Extractor (2/2)

- Example: perform feature extraction given a list

```
print("Create Feature extractor")
extractor = sidekit.FeaturesExtractor(audio_filename_structure=None,
                                     feature_filename_structure=None,
                                     sampling_frequency=None,
                                     lower_frequency=200,
                                     higher_frequency=3800,
                                     filter_bank="log",
                                     filter_bank_size=24,
                                     window_size=0.025,
                                     shift=0.01,
                                     ceps_number=20,
                                     vad="snr",
                                     snr=40,
                                     pre_emphasis=0.97,
                                     save_param=["vad", "energy", "cep", "fb"],
                                     keep_all_features=True)
```

Initiate a FeatureExtractor
with **attributes** given as
arguments

```
print("Start extracting features")
extractor.save_list(show_list=show_list,
                   channel_list=channel_list,
                   audio_file_list=audio_list,
                   feature_file_list=feature_list,
                   num_thread=nbThread)
```

```
show_list[0] => 'sre04/tghy'
channel_list[0] => 0
audio_list[0] => '/media/kalee/SRE04/train/data/tghy.sph'
feature_list[0] => '/disk1/ProjectPy/sidekit3/mfcc_24/sre04/tghy.h5'
```

Extract feature from
channel a or b from a
audio file and save all
features to a hdf5 file

Feature Server (1/2)

- FeatureServer – a Python class providing functionality to load features from files and post-process the features:
 - Normalization (CMS, CMVN, RASTA, Short-Term Gaussianization)
 - Temporal context (delta, delta-delta)
 - Feature selection
- Initiate a FeatureServer

```
# Define a FeaturesServer for loading the acoustic features
fs = sidekit.FeaturesServer(feature_filename_structure="{dir}/{}/{}.{ext}".format(dir=
feature_dir, ext=feature_extension),
```

'h5' ↗

```
    dataset_list=["energy", "cep", "vad"],
    mask="[0-12]",
    feat_norm="cmvn",
    keep_all_features=False,
    delta=True,
    double_delta=True,
    rasta=True,
    context=None)
```

Perform RASTA,
append delta and
delta-delta

Load and concatenate
cepstral coefficients (first
13 from 0 to 12) and log
energy

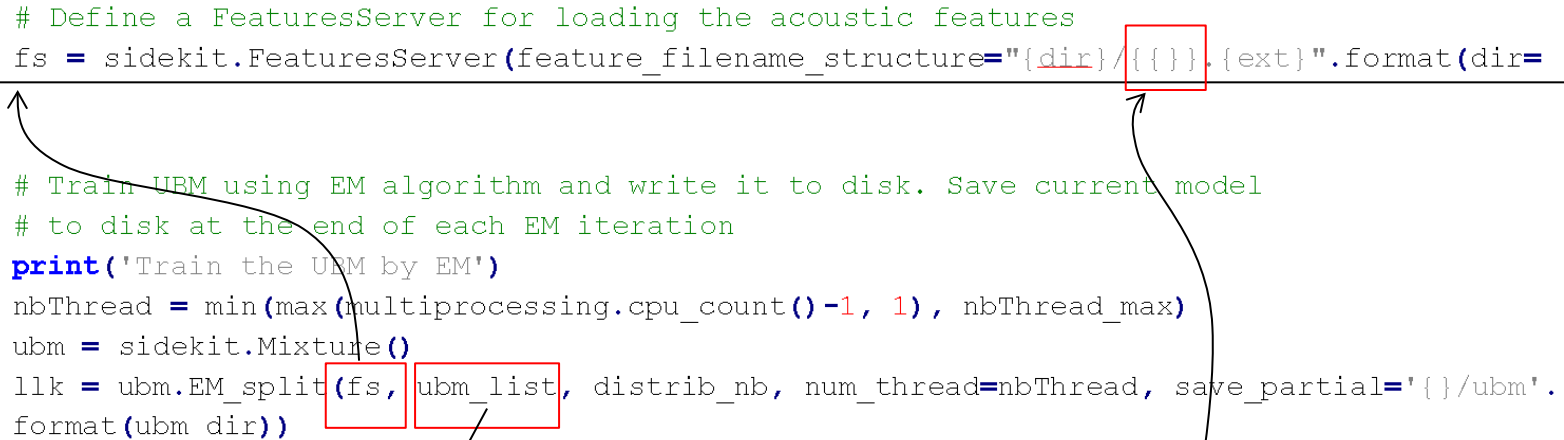
Remove silence frames
based on the VAD labels

Feature Server (2/2)

- Example: Using FeatureServer to load features for UBM training

```
# Define a FeaturesServer for loading the acoustic features
fs = sidekit.FeaturesServer(feature_filename_structure="{dir}/{[]}.ext".format(dir=

# Train UBM using EM algorithm and write it to disk. Save current model
# to disk at the end of each EM iteration
print('Train the UBM by EM')
nbThread = min(max(multiprocessing.cpu_count()-1, 1), nbThread_max)
ubm = sidekit.Mixture()
llk = ubm.EM_split(fs, ubm_list, distrib_nb, num_thread=nbThread, save_partial='{}/ubm'.
format(ubm_dir))
```



List of feature files to train UBM

```
sre06/noil_b
sre06/pbgy_a
sre05/jbwz_a
.
.
```

Feature filename
and the dataset
name in the h5 file.

Stat Server (1/2)

- StatServer – a Python class providing functionality to extract, store, process, train total variability model (or PLDA) and extract i-vector.
- Extract sufficient statistics

```
back_stat = sidekit.StatServer(back_idmap, ubm)
back_stat.accumulate_stat(ubm=ubm, feature_server=fs, seg_indices=range(back_stat.segset.shape[0]), num_threads=num_threads)
back_stat.write('data/stat_back_{}.h5'.format(distrib_nb))
```

- Train TVM (with sufficient statistics whitening and MD re-estimation)

```
tv_stat = sidekit.StatServer.read_subset('data/stat_back_{}.h5'.format(distrib_nb), tv_idmap)
tv_mean, tv, __, __, tv_sigma = tv_stat.factor_analysis(rank_f = rank_TV,
                                                         rank_g = 0,
                                                         rank_h = None,
                                                         re_estimate_residual = False,
                                                         it_nb = (tv_iteration, 0, 0),
                                                         min_div = True,
                                                         ubm = ubm,
                                                         batch_size = 100,
                                                         num_thread = nbThread,
                                                         save_partial = "data/TV_{}".format(distrib_nb))
sidekit.sidekit_io.write_tv_hdf5((tv, tv_mean, tv_sigma), "data/TV_{}".format(distrib_nb))
```

Stat Server (2/2)

- Extract sufficient statistics

```
enroll_stat = sidekit.StatServer(enroll_idmap, ubm)
enroll_stat.accumulate_stat(ubm=ubm, feature_server=fs, seg_indices=range(enroll_stat.segset.shape[0]), num_
enroll_stat.write('data/stat_sre10_core-core_enroll_{}.h5'.format(distrib_nb))

test_stat = sidekit.StatServer(test_idmap, ubm)
test_stat.accumulate_stat(ubm=ubm, feature_server=fs, seg_indices=range(test_stat.segset.shape[0]), num_thre
test_stat.write('data/stat_sre10_core-core_test_{}.h5'.format(distrib_nb))

back_idmap = plda_all_idmap.merge(tv_idmap)
back_stat = sidekit.StatServer(back_idmap, ubm)
back_stat.accumulate_stat(ubm=ubm, feature_server=fs, seg_indices=range(back_stat.segset.shape[0]), num_thre
back_stat.write('data/stat_back_{}.h5'.format(distrib_nb))
```

- Extract i-vector for Train, Enroll and Test i-vectors

```
enroll_stat = sidekit.StatServer('data/stat_sre10_core-core_enroll_{}.h5'.format(distrib_nb))
enroll_iv = enroll_stat.estimate_hidden(tv_mean, tv_sigma, V=tv, batch_size=100, num_thread=nbThread)[0]
enroll_iv.write('data/iv_sre10_core-core_enroll_{}.h5'.format(distrib_nb))

test_stat = sidekit.StatServer('data/stat_sre10_core-core_test_{}.h5'.format(distrib_nb))
test_iv = test_stat.estimate_hidden(tv_mean, tv_sigma, V=tv, batch_size=100, num_thread=nbThread)[0]
test_iv.write('data/iv_sre10_core-core_test_{}.h5'.format(distrib_nb))

plda_stat = sidekit.StatServer.read_subset('data/stat_back_{}.h5'.format(distrib_nb), plda_all_idmap)
plda_iv = plda_stat.estimate_hidden(tv_mean, tv_sigma, V=tv, batch_size=100, num_thread=nbThread)[0]
plda_iv.write('data/iv_plda_{}.h5'.format(distrib_nb))
```


Whitening (1/5)

- Whitening or sphering [Bishop, 2006] refers to the normalization of data, where $\boldsymbol{\mu}$ and $\boldsymbol{\Phi}$ are sample mean and covariance matrix estimated from data:

$$\tilde{o}_t = \boldsymbol{\Phi}^{-1/2}(\mathbf{o}_t - \boldsymbol{\mu})$$

$$\boldsymbol{\mu} = \frac{1}{T} \sum_{t=1}^T \mathbf{o}_t, \quad \boldsymbol{\Phi} = \frac{1}{T} \sum_{t=1}^T (\mathbf{o}_t - \boldsymbol{\mu})(\mathbf{o}_t - \boldsymbol{\mu})^T$$

- The matrix $\boldsymbol{\Phi}^{-1/2}$ is obtained by decomposing $\boldsymbol{\Phi}$ to two parts, for instance, with Cholesky decomposition or eigenvalue decomposition. In both cases, a better choice is to decompose $\boldsymbol{\Phi}$ and then followed by inversion
 - Cholesky decomposition:

$$\boldsymbol{\Phi}^{-1} = (\mathbf{L}\mathbf{L}^T)^{-1} = (\mathbf{L}^T)^{-1}(\mathbf{L}^{-1}) = (\mathbf{L}^{-1})^T(\mathbf{L}^{-1})$$

$$\tilde{o}_t = \mathbf{L}^{-1}(\mathbf{o}_t - \boldsymbol{\mu})$$


Whitening (2/5)

- Eigenvalue decomposition:

$$\Phi^{-1} = (\mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T)^{-1} = \mathbf{Q}\mathbf{\Lambda}^{-1}\mathbf{Q}^T = \mathbf{Q}\mathbf{\Lambda}^{-1/2}\mathbf{\Lambda}^{-1/2}\mathbf{Q}^T = (\mathbf{Q}\mathbf{\Lambda}^{-1/2})(\mathbf{Q}\mathbf{\Lambda}^{-1/2})^T$$

$$\tilde{o}_t = (\mathbf{Q}\mathbf{\Lambda}^{-1/2})^T(o_t - \boldsymbol{\mu}) = \mathbf{\Lambda}^{-1/2}\mathbf{Q}^T(o_t - \boldsymbol{\mu})$$

- Whitening leads to zero mean and unit covariance distribution



The effects of whitening: original data (left), whitening with diagonal covariance matrix (middle), whitening with full covariance matrix (right) [Bishop, 2006]

Whitening (3/5)

- We aim to whiten the acoustic observations with respect to the UBM (multiple Gaussians). This is accomplished by subjecting the first-order statistics to the following transformation.
 - Soft-alignment of frames to component

$$\mathbf{F}_r(c) = \sum_t \gamma_t(c) o_t$$

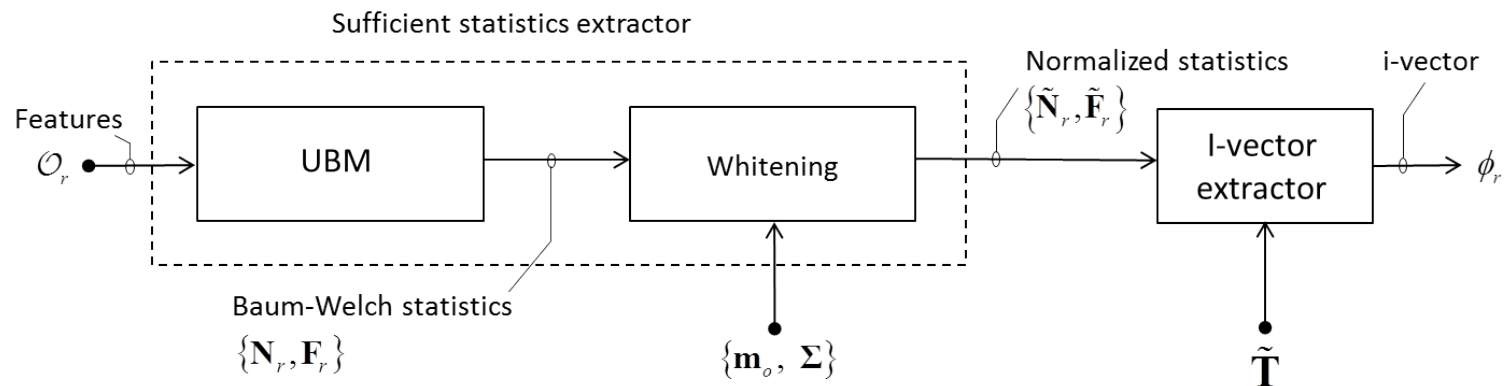
- Subjecting the first-order statistics to the following transformation

$$\tilde{\mathbf{F}}_r(c) = \boldsymbol{\Phi}_c^{-1/2} \sum_t \gamma_t(c) (o_t - \boldsymbol{\mu}_c) \text{ for } c = 1, 2, \dots, C$$

$$\tilde{\mathbf{F}}_r = \boldsymbol{\Sigma}^{-1/2} (\mathbf{F}_r - \mathbf{N}_r \mathbf{m}_o)$$

Whitening (4/5)

- UBM parameters $\{\mathbf{m}_o, \mathbf{\Sigma}\}$ are absorbed as part of sufficient statistics and therefore no longer required in subsequent processing (as shown in the block diagram).
- The implementation becomes the same for diagonal or full covariance matrices.
- Strictly speaking, an i-vector extractor consists both sufficient statistics extraction and i-vector inference.



Whitening (5/5)

- Using the whitened sufficient statistics for training, we obtained a whitened total variability matrix $\tilde{\mathbf{T}}$ (i.e., whitening the sufficient statistics has similar effects in whitening the \mathbf{T} matrix). The original \mathbf{T} matrix could be recovered by inverting the transformation

$$\tilde{\mathbf{T}} = \mathbf{\Sigma}^{-1/2} \mathbf{T} \quad \Rightarrow \quad \mathbf{T} = \mathbf{\Sigma}^{1/2} \tilde{\mathbf{T}}$$

- The i-vectors (more generally the posterior distribution of the latent variable) remains the same

$$\begin{aligned} \phi_r &= \left(\mathbf{I} + \mathbf{T}^T \mathbf{\Sigma}^{-1} \mathbf{N}_r \mathbf{T} \right)^{-1} \cdot \mathbf{T}^T \mathbf{\Sigma}^{-1} (\mathbf{F}_r - \mathbf{N}_r \mathbf{m}_o) \\ &= \left(\mathbf{I} + \mathbf{T}^T \mathbf{\Sigma}^{-1/2} \mathbf{\Sigma}^{-1/2} \mathbf{N}_r \mathbf{T} \right)^{-1} \cdot \mathbf{T}^T \mathbf{\Sigma}^{-1/2} \mathbf{\Sigma}^{-1/2} (\mathbf{F}_r - \mathbf{N}_r \mathbf{m}_o) \\ &= \left[\mathbf{I} + \left(\mathbf{\Sigma}^{-1/2} \mathbf{T} \right)^T \mathbf{N}_r \left(\mathbf{\Sigma}^{-1/2} \mathbf{T} \right) \right]^{-1} \cdot \left(\mathbf{\Sigma}^{-1/2} \mathbf{T} \right)^T \mathbf{\Sigma}^{-1/2} (\mathbf{F}_r - \mathbf{N}_r \mathbf{m}_o) \\ &= \underbrace{\left[\mathbf{I} + \tilde{\mathbf{T}}^T \mathbf{N}_r \tilde{\mathbf{T}} \right]^{-1}}_{\mathbf{L}_r^{-1}} \cdot \tilde{\mathbf{T}}^T \tilde{\mathbf{F}}_r \end{aligned}$$

Minimum divergence re-estimation (1/3)

- The objective of minimum divergence re-estimation is to scale the \mathbf{T} matrix (and shift the global mean vector \mathbf{m}_o) in order to force the empirical distribution of the i-vectors conform to the standard Gaussian prior.
- Consider a TVM with single Gaussian and single frame for the purpose of illustration

$$\begin{aligned} p(o) &= \int \mathcal{N}(o | \mathbf{m}_o + \mathbf{T}\mathbf{h}, \Phi) \mathcal{N}(\mathbf{h} | \boldsymbol{\mu}_h, \boldsymbol{\Sigma}_h) d\mathbf{h} \\ &= \mathcal{N}(o | \mathbf{m}_o + \mathbf{T}\boldsymbol{\mu}_h, \mathbf{T}\boldsymbol{\Sigma}_h\mathbf{T}^T + \Phi) \\ &= \mathcal{N}(o | \mathbf{m}_o + \mathbf{T}\boldsymbol{\mu}_h, \mathbf{T}(\mathbf{L}\mathbf{L}^T)\mathbf{T}^T + \Phi) \\ &= \mathcal{N}(o | \mathbf{m}_o + \mathbf{T}(\mathbf{L}\mathbf{L}^{-1})\boldsymbol{\mu}_h, (\mathbf{T}\mathbf{L})(\mathbf{T}\mathbf{L})^T + \Phi) \\ &= \mathcal{N}(o | \mathbf{m}_o + \tilde{\mathbf{T}}\mathbf{L}^{-1}\boldsymbol{\mu}_h, \tilde{\mathbf{T}}\tilde{\mathbf{T}}^T + \Phi) \\ &= \int \mathcal{N}(o | \mathbf{m}_o + \tilde{\mathbf{T}}\mathbf{L}^{-1}\boldsymbol{\mu}_h + \tilde{\mathbf{T}}\mathbf{h}, \Phi) \mathcal{N}(\mathbf{h} | 0, \mathbf{I}) d\mathbf{h} \end{aligned}$$

Minimum divergence re-estimation (2/3)

- MD re-estimation of \mathbf{T} matrix and global mean vector \mathbf{m}_o

$$\begin{aligned}\tilde{\mathbf{T}} &= \mathbf{T}\mathbf{L} \\ \tilde{\mathbf{m}}_o &= \mathbf{m}_o + \tilde{\mathbf{T}}\mathbf{L}^{-1}\boldsymbol{\mu}_h = \mathbf{m}_o + \mathbf{T}\boldsymbol{\mu}_h\end{aligned}$$

- In general, most implementations rotate and scale only the \mathbf{T} matrix

$$\begin{aligned}p(o) &= \int \mathcal{N}(o | \mathbf{m}_o + \mathbf{T}\mathbf{h}, \boldsymbol{\Phi}) \mathcal{N}(\mathbf{h} | \boldsymbol{\mu}_h, \boldsymbol{\Sigma}_h) d\mathbf{h} \\ &= \mathcal{N}\left(o | \mathbf{m}_o + \tilde{\mathbf{T}}\mathbf{L}^{-1}\boldsymbol{\mu}_h, \tilde{\mathbf{T}}\tilde{\mathbf{T}}^T + \boldsymbol{\Phi}\right) \\ &= \int \mathcal{N}(o | \mathbf{m}_o + \tilde{\mathbf{T}}\mathbf{h}, \boldsymbol{\Phi}) \mathcal{N}(\mathbf{h} | \mathbf{L}^{-1}\boldsymbol{\mu}_h, \mathbf{I}) d\mathbf{h}\end{aligned}$$

- MD re-estimation is applied after each EM step. Its role is to get good estimate of the eigenvalues corresponding to the eigenvectors defining the total variability space [Kenny, 2008]

Minimum divergence re-estimation (3/3)

- The empirical distribution $N(\mathbf{h}|\boldsymbol{\mu}_h, \boldsymbol{\Sigma}_h)$ is estimated from i-vectors minimizing the KL divergence criterion, i.e., to find the Gaussian distribution $N(\mathbf{h}|\boldsymbol{\mu}_h, \boldsymbol{\Sigma}_h)$ with minimum KL distances from all the R posterior distributions

$$D(\theta_{\text{MD}}) = \sum_{i=1}^I E \left\{ \log \frac{\mathcal{N}(\mathbf{h} | \phi_i, \mathbf{L}_i^{-1})}{\mathcal{N}(\mathbf{h} | \boldsymbol{\mu}_h, \boldsymbol{\Sigma}_h)} \right\}$$

- Closed-form solution

$$\boldsymbol{\mu}_h = \frac{1}{R} \sum_{r=1}^R \phi_r$$

$$\boldsymbol{\Sigma}_h = \left(\frac{1}{R} \sum_{r=1}^R L_r^{-1} + \phi_r \phi_r^T \right) - \boldsymbol{\mu}_h \boldsymbol{\mu}_h^T$$

Appendix – MD estimate of prior (1/4)

$$\begin{aligned}
 D(\theta) &= \sum_{r=1}^R D_{KL}(N_r | N_0) \quad \leftarrow \text{number of sessions/posterior distributions} \\
 &= \sum_{r=1}^R E \left\{ \log \frac{N(h|m_r, L_r^{-1})}{N(h|y, P^{-1})} \right\}, \quad \theta = \{y, P^{-1}\} \quad \begin{matrix} \mu_h & \Sigma_h \\ \downarrow & \downarrow \end{matrix} \\
 &= \sum_{r=1}^R \frac{1}{2} \left[\text{tr}(P L_r^{-1}) + (y - m_r)^T P (y - m_r) \right. \\
 &\quad \left. - k - \log |L_r^{-1}| + \log |P^{-1}| \right] \\
 &\quad \uparrow \\
 &\quad \text{Dimensionality of variables} \\
 &= \sum_{r=1}^R \frac{1}{2} \left[\text{tr}(P L_r^{-1}) + (y - m_r)^T P (y - m_r) \right. \\
 &\quad \left. + \log |P^{-1}| - k - \log |L_r^{-1}| \right] \\
 &\quad \underbrace{\hspace{10em}}_{\text{constant for a given dimensionality and dataset}}
 \end{aligned}$$

Appendix – MD estimate of prior (2/4)

$$D(\theta) = \frac{1}{2} \left[\sum_{r=1}^R \text{tr}(L_r^{-1} P) + \underbrace{(y - m_r)^T P (y - m_r)}_{\text{tr}((y - m_r)(y - m_r)^T P)} \right] - \frac{1}{2} R \log |P| + C \quad (1)$$

Differentiate (1) wrt y

$$\frac{dD(\theta)}{dy} = \sum_{r=1}^R (y - m_r)^T P \quad (2)$$

Setting (2) to zero

$$\sum_{r=1}^R (y - m_r)^T P = 0$$

$$\sum_{r=1}^R (y - m_r) = 0$$

$$Ry - \sum_{r=1}^R m_r = 0$$

$$y = \frac{1}{R} \sum_{r=1}^R m_r \quad (4)$$

Appendix – MD estimate of prior (3/4)

$$D(\theta) = \frac{1}{2} \left[\text{tr}(SP) - R \log |P| \right] + C \quad \leftarrow \text{②}$$

$$S = \sum_{r=1}^R L_r^{-1} + (y - m_r)(y - m_r)^T$$

Differentiate (2) wrt P

$$\frac{dD(\theta)}{dP} = \frac{1}{2}(S - RP^{-1}) \quad \text{--- ⑤}$$

setting (5) to zero

$$\begin{aligned} S &= RP^{-1} \\ P^{-1} &= \frac{1}{R} S = \frac{1}{R} \sum_{r=1}^R L_r^{-1} + (y - m_r)(y - m_r)^T \\ &= L^{-1} + \frac{1}{R} \sum_{r=1}^R (y - m_r)(y - m_r)^T \end{aligned}$$

Appendix – MD estimate of prior (4/4)

$$\begin{aligned} P^{-1} &= \frac{1}{R} \sum_{r=1}^R L_r^{-1} + \overbrace{(y - m_r)(y - m_r)^T}^{\text{Empirical term}} \\ &= \frac{1}{R} \sum_{r=1}^R L_r^{-1} + y y^T + m_r m_r^T - 2 y m_r^T \\ &= \frac{1}{R} \sum_{r=1}^R L_r^{-1} + m_r m_r^T + \frac{1}{R} \sum_{r=1}^R (y y^T - 2 y m_r^T) \\ &= \frac{1}{R} \sum_{r=1}^R L_r^{-1} + m_r m_r^T + y y^T - 2 y y^T \\ &= \left(\frac{1}{R} \sum_{r=1}^R L_r^{-1} + m_r m_r^T \right) - \underbrace{y y^T}_{\mu_h \mu_h^T} \end{aligned}$$

SIDEKIT – Python Toolkit for Speaker Recognition

PROBABILISTIC LDA (PLDA)

Channel compensation

- I-vector extraction is a **compression** process, where we compress across the time (variable-length sequence to **fixed-length** vector) and supervector space (high to **low** dimensionality).
- An i-vector contains
 - Speaker/language characteristic
 - other information originated from the transmission channel, recording devices, and acoustic environment (**channel** or **session** variability)
- Channel (or session) compensation – to deal with the **mismatch** between enrolment and test utterances.
- Main stream compensation techniques include (implemented in SIDEKIT)
 - Linear discriminant analysis (LDA) [Bishop, 2006]
 - Within-class covariance normalization [Hatch et al, 2006]
 - **Probabilistic LDA** (PLDA) [Prince and Elder, 2007]

I-vector/PLDA using SIDEKIT

- PLDA scoring

```
# Normalize i-vector using Spherical Nuisance Normalization (similar to  
length-normalization) and compute scores using PLDA
```

```
meanSN, CovSN = plda_iv.estimate_spectral_norm_stat1(1, 'sphNorm')
```

```
plda_iv.spectral_norm_stat1(meanSN[:1], CovSN[:1])
```

```
enroll_iv.spectral_norm_stat1(meanSN[:1], CovSN[:1])
```

```
test_iv.spectral_norm_stat1(meanSN[:1], CovSN[:1])
```

Whitening + unit norm

```
plda_mean, plda_F, plda_G, plda_H, plda_Sigma = plda_iv.factor_analysis(rank_f=plda_rk,
```

```
rank_g=0,
```

```
rank_h=None,
```

```
re_estimate_residual=  
True,
```

```
it_nb=(10,0,0),
```

```
min_div=True,
```

```
ubm=None,
```

```
batch_size=1000,
```

```
num_thread=nbThread)
```

NOT used in simplified PLDA

Channel variability is model
with full residual covariance
matrix for simplified PLDA

```
scores_plda = sidekit.iv_scoring.PLDA_scoring(enroll_iv, test_iv, test_ndx, plda_mean, plda_F  
, plda_G, plda_Sigma, full_model=False)
```

Cosine scoring

- LDA followed by cosine scoring

```
# LDA (reduce the dimension of i-vectors to 150 dimensions) followed by cosine scoring.
LDA = plda_iv.get_lda_matrix_stat1(150)
plda_iv.rotate_stat1(LDA)
enroll_iv.rotate_stat1(LDA)
test_iv.rotate_stat1(LDA)
scores_cos_lda = sidekit.iv_scoring.cosine_scoring(enroll_iv, test_iv, test_ndx, wccn=None)
```

- WCCN followed by cosine scoring

```
# WCCN followed by cosine scoring.
wccn = plda_iv.get_wccn_choleski_stat1()
scores_cos_wccn = sidekit.iv_scoring.cosine_scoring(enroll_iv, test_iv, test_ndx, wccn=wccn)
```


PLDA definition

- PLDA explains the observed data (i.e., i-vectors) in terms of **speaker** identity and **channel** effects.
- Let ϕ_{ij} be an i-vector extracted from the j -th session of speaker i , the generative explanation for the i-vector is

$$\phi_{ij} = \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i + \mathbf{G}\mathbf{w}_{ij} + \boldsymbol{\varepsilon}_{ij} \text{ for } i = 1, 2, \dots, I \text{ and } j = 1, 2, \dots, J$$

Number of speakers Number of sessions per speaker

↓ ↓

- The **speaker** variable \mathbf{h}_i quantifies the observed deviations from the mean $\boldsymbol{\mu}$ due to the changes of speaker.
- The **channel** variables \mathbf{w}_{ij} quantifies the observed deviations from the mean $\boldsymbol{\mu}$ due to the different sessions of the same speaker
- The **residual** noise term $\boldsymbol{\varepsilon}_{ij} \sim N(0, \boldsymbol{\Phi})$ described the residual variation (note: this corresponds to the UBM covariance matrices in the total variability model)

Probabilistic representation

- In probabilistic term, PLDA is expressed as

$$p(\phi_{ij} | \mathbf{h}_i, \mathbf{w}_{ij}) = \mathcal{N}(\phi_{ij} | \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i + \mathbf{G}\mathbf{w}_{ij}, \boldsymbol{\Phi})$$

Conditional distribution

Priors $\left\{ \begin{array}{l} p(\mathbf{h}_i) = \mathcal{N}(\mathbf{h}_i | 0, \mathbf{I}) \\ p(\mathbf{w}_{ij}) = \mathcal{N}(\mathbf{w}_{ij} | 0, \mathbf{I}) \end{array} \right.$

- Given the **prior** and **conditional** distributions, we integrate out the speaker and channel latent variables to obtain the **marginal** distribution.

$$p(\phi_{ij}) = \int p(\phi_{ij} | \mathbf{h}_i, \mathbf{w}_{ij}) p(\mathbf{h}_i) p(\mathbf{w}_{ij}) d\mathbf{h}_i d\mathbf{w}_{ij}$$
$$= \mathcal{N}(\phi_{ij} | \boldsymbol{\mu}, \mathbf{F}\mathbf{F}^T + \mathbf{G}\mathbf{G}^T + \boldsymbol{\Phi})$$

Structured covariance matrix


Between speaker variation \rightarrow $\mathbf{F}\mathbf{F}^T$

Within speaker variation \rightarrow $\mathbf{G}\mathbf{G}^T$

- PLDA is a Gaussian density with **structured covariance matrix**, similar to factor analysis, with additional subspaces **F** and **G**.

Graphical model representation

- A PLDA is a **factor analysis** model with additional **tying across observations**
 - I-vectors are the observed variables
 - I-vectors of the same speaker share the speaker variable \mathbf{h}_i (i.e., tying of latent variable across i-vectors of the same speaker)
 - Each i-vector is characterized by a channel variable \mathbf{w}_{ij}
 - Separate speaker and channel subspaces $\{\mathbf{F}, \mathbf{G}\}$ to tease apart the contributions of speaker and channel information (TVM has one single subspace)



Simplified PLDA

- The general PLDA is a model for Gaussian distribution with a structured covariance matrix

$$\phi_{ij} = \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i + \mathbf{G}\mathbf{w}_{ij} + \boldsymbol{\varepsilon}_{ij} \text{ where } \boldsymbol{\varepsilon}_{ij} \sim \mathcal{N}(0, \boldsymbol{\Phi})$$

$$p(\phi_{ij} | \mathbf{h}_i, \mathbf{w}_{ij}) = \mathcal{N}(\phi_{ij} | \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i + \mathbf{G}\mathbf{w}_{ij}, \boldsymbol{\Phi})$$

$$p(\phi_{ij}) = \mathcal{N}(\phi_{ij} | \boldsymbol{\mu}, \mathbf{F}\mathbf{F}^T + \mathbf{G}\mathbf{G}^T + \boldsymbol{\Phi})$$


- The simplified PLDA is obtained by integrate out \mathbf{w}

$$\phi_{ij} = \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i + \tilde{\boldsymbol{\varepsilon}}_{ij} \text{ where } \tilde{\boldsymbol{\varepsilon}}_{ij} \sim \mathcal{N}(0, \tilde{\boldsymbol{\Phi}})$$

$$p(\phi_{ij} | \mathbf{h}_i) = \mathcal{N}(\phi_{ij} | \boldsymbol{\mu} + \mathbf{F}\mathbf{h}_i, \tilde{\boldsymbol{\Phi}})$$

$$p(\phi_{ij}) = \mathcal{N}(\phi_{ij} | \boldsymbol{\mu}, \mathbf{F}\mathbf{F}^T + \underbrace{\mathbf{G}\mathbf{G}^T + \boldsymbol{\Phi}}_{\tilde{\boldsymbol{\Phi}}})$$

Full covariance matrix



Posterior inference and EM

- Given all the J i-vectors from speaker i , the posterior distribution of \mathbf{h}_i is a normal distribution

$$p(\mathbf{h}_i | \phi_{i,1}, \phi_{i,2}, \dots, \phi_{i,J}) = \mathcal{N}(\mathbf{h}_i | \mathbf{m}_h, \mathbf{L}_h^{-1})$$

Posterior mean

Posterior covariance

$$\mathbf{m}_h = \underbrace{(\mathbf{J}\mathbf{F}^T\tilde{\Phi}^{-1}\mathbf{F} + \mathbf{I})^{-1}}_{\mathbf{L}_h^{-1}} \cdot \mathbf{F}^T\tilde{\Phi}^{-1} \left(\sum_{j=1}^J \phi_{ij} - \boldsymbol{\mu} \right)$$

- The parameters of the simplified PLDA model $\theta = \{\mathbf{F}, \tilde{\Phi}\}$ are learned via EM algorithm
 - E-step computes the posterior distribution for each speaker variables \mathbf{h}_i
 - M-step updates the using the relations

$$\mathbf{F} = \underbrace{\left(\sum_{i=1}^I \sum_{j=1}^J (\phi_{ij} - \boldsymbol{\mu}) \mathbf{E}[\mathbf{h}_i^T] \right)}_{\text{First-order statistics}} \underbrace{\left(\sum_{i=1}^I \mathbf{J} \mathbf{E}[\mathbf{h}_i \mathbf{h}_i^T] \right)^{-1}}_{\text{Zero-order statistics}}$$


$$\tilde{\Phi} = \sum_{i=1}^I \sum_{j=1}^J (\phi_{ij} - \boldsymbol{\mu})(\phi_{ij} - \boldsymbol{\mu})^T - \mathbf{F} \mathbf{E}[\mathbf{h}_i] (\phi_{ij} - \boldsymbol{\mu})^T$$

PLDA scoring

- A commonly used scoring method is based on the likelihood ratio test between two hypotheses – whether **the enrollment and test i-vectors, ϕ_e and ϕ_t , are from the same speaker or different speakers.**
- The verification score is given by the log-likelihood ratio between the two models $\{M_0, M_1\}$

$$l(\phi_e, \phi_t) = \log \frac{p(\phi_e, \phi_t | M_0)}{p(\phi_e, \phi_t | M_1)}$$

$$= \log \frac{\mathcal{N}\left(\begin{bmatrix} \phi_e \\ \phi_t \end{bmatrix} \middle| \begin{bmatrix} \mu \\ \mu \end{bmatrix}, \begin{bmatrix} \mathbf{F}\mathbf{F}^T + \tilde{\Phi} & \mathbf{F}\mathbf{F}^T \\ \mathbf{F}\mathbf{F}^T & \mathbf{F}\mathbf{F}^T + \tilde{\Phi} \end{bmatrix}\right)}{\mathcal{N}(\phi_e | \mu, \mathbf{F}\mathbf{F}^T + \tilde{\Phi}) \mathcal{N}(\phi_t | \mu, \mathbf{F}\mathbf{F}^T + \tilde{\Phi})}$$



- We don't build "speaker model". Speaker verification scores are computed by comparing the enrollment and test i-vectors through PLDA model.

Appendix: joint distribution $p(\phi_e, \phi_t)$

- Construct the composite equation

$$\underbrace{\begin{bmatrix} \phi_e \\ \phi_t \end{bmatrix}}_{\phi'} = \underbrace{\begin{bmatrix} \boldsymbol{\mu} \\ \boldsymbol{\mu} \end{bmatrix}}_{\boldsymbol{\mu}'} + \underbrace{\begin{bmatrix} \mathbf{F} \\ \mathbf{F} \end{bmatrix}}_{\mathbf{F}'} \mathbf{h} + \underbrace{\begin{bmatrix} \boldsymbol{\varepsilon}_e \\ \boldsymbol{\varepsilon}_t \end{bmatrix}}_{\boldsymbol{\varepsilon}'}$$
$$\phi' = \boldsymbol{\mu}' + \mathbf{F}'\mathbf{h} + \boldsymbol{\varepsilon}'$$

- Use the composite terms in the PLDA marginal distribution

$$\begin{aligned} p(\phi') &= \int \mathcal{N}(\phi' | \boldsymbol{\mu}' + \mathbf{F}'\mathbf{h}, \boldsymbol{\Phi}') \mathcal{N}(\mathbf{h} | 0, \mathbf{I}) d\mathbf{h} \\ &= \mathcal{N}(\phi' | \boldsymbol{\mu}', \mathbf{F}'\mathbf{F}'^T + \boldsymbol{\Phi}') \\ \mathbf{F}'\mathbf{F}'^T + \boldsymbol{\Phi}' &= \begin{bmatrix} \mathbf{F} \\ \mathbf{F} \end{bmatrix} \begin{bmatrix} \mathbf{F}^T & \mathbf{F}^T \end{bmatrix} + \begin{bmatrix} \tilde{\boldsymbol{\Phi}} & 0 \\ 0 & \tilde{\boldsymbol{\Phi}} \end{bmatrix} \\ &= \begin{bmatrix} \mathbf{F}\mathbf{F}^T + \tilde{\boldsymbol{\Phi}} & \mathbf{F}\mathbf{F}^T \\ \mathbf{F}\mathbf{F}^T & \mathbf{F}\mathbf{F}^T + \tilde{\boldsymbol{\Phi}} \end{bmatrix} \end{aligned}$$

Speaker adaptation (1/2)

- Re-write the log-likelihood ratio as follows

$$l(\phi_e, \phi_t) = \log \frac{p(\phi_e, \phi_t | M_0)}{p(\phi_e, \phi_t | M_1)} = \log \frac{p(\phi_t | \phi_e) \cancel{p(\phi_e)}}{p(\phi_t) \cancel{p(\phi_e)}} = \log \frac{p(\phi_t | \phi_e)}{p(\phi_t)}$$

- The numerator and denominator could be interpreted as **speaker-dependent and universal PLDA models**

$$\begin{aligned} p(\phi_t | \phi_e) &= \int p(\phi_t | \mathbf{h}) p(\mathbf{h} | \phi_e) d\mathbf{h} \\ &= \mathcal{N}\left(\phi_t \mid \underbrace{\boldsymbol{\mu} + \mathbf{F}\mathbf{m}_h}_{\text{Adapted mean vector}}, \underbrace{\mathbf{F}\mathbf{L}_h^{-1}\mathbf{F}^T + \tilde{\boldsymbol{\Phi}}}_{\text{Adapted covariance matrix}}\right) \end{aligned} \quad \left| \quad \begin{aligned} p(\phi_t) &= \int p(\phi_t | \mathbf{h}) p(\mathbf{h}) d\mathbf{h} \\ &= \mathcal{N}\left(\phi_t \mid \boldsymbol{\mu}, \mathbf{F}\mathbf{F}^T + \tilde{\boldsymbol{\Phi}}\right) \end{aligned}$$

Adapted mean vector

Adapted covariance matrix

- The verification score is given by the log-likelihood ratio between the speaker-dependent PLDA model and the universal PLDA model, in a way much similar to the idea of *Universal Background Model*.


Speaker adaptation (2/2)

- The parameters $\{\mathbf{m}_h, \mathbf{L}_h^{-1}\}$ are the posterior mean and covariance of the latent speaker variable estimated using the enrollment i-vector ϕ_e .

$$p(\mathbf{h} | \phi_e) = \mathcal{N}(\mathbf{h} | \mathbf{m}_h, \mathbf{L}_h^{-1})$$
$$\mathbf{m}_h = \mathbf{L}_h^{-1} \cdot \mathbf{F}^T \tilde{\Phi}^{-1} (\phi_e - \boldsymbol{\mu})$$
$$\mathbf{L}_h^{-1} = (\mathbf{F}^T \tilde{\Phi}^{-1} \mathbf{F} + I)^{-1}$$

- Extension to multi-session enrollment (by-the-book solution)

$$\mathbf{m}_h = \mathbf{L}_h^{-1} \cdot \mathbf{F}^T \tilde{\Phi}^{-1} \left(\sum_{j=1}^J \phi_{ij} - \boldsymbol{\mu} \right)$$
$$\mathbf{L}_h^{-1} = \left(J \mathbf{F}^T \tilde{\Phi}^{-1} \mathbf{F} + I \right)^{-1} \longleftarrow \text{The covariance shrinks for larger } J$$



Graphical model illustrating the model adaptation scoring approach

L. Chen, K. A. Lee et al, "Minimum divergence estimation of speaker prior in multi-session PLDA scoring," in *Proc. IEEE ICASSP*, 2014, pp. 4035 – 4039.

Minimum divergence estimate of speaker prior (1/2)

- For each enrollment session from the speaker i , we compute the mean and covariance of the posterior distribution:

$$\mathbf{m}_{ij} = \mathbf{L}^{-1} \mathbf{F}^T \tilde{\mathbf{\Phi}}^{-1} (\phi_{ij} - \boldsymbol{\mu})$$

$$\mathbf{L}^{-1} = [\mathbf{F}^T \tilde{\mathbf{\Phi}}^{-1} \mathbf{F} + \mathbf{I}]^{-1} \longleftarrow \text{The same for all sessions}$$

- We seek for a Gaussian distribution (i.e., the speaker prior) $N(\mathbf{h}|\mathbf{y}_i, \mathbf{P}_i^{-1})$ that best represents the J posterior distributions [Chen et al, ICASSP 2014] minimizing the KL divergence between the prior from the J posteriors:

$$D(\theta_{\text{MD}}) = \sum_{j=1}^J E \left\{ \log \frac{\mathcal{N}(\mathbf{h} | \mathbf{m}_{ij}, \mathbf{L}^{-1})}{\mathcal{N}(\mathbf{h} | \mathbf{y}_i, \mathbf{P}_i^{-1})} \right\}$$

$$\mathcal{N}(\mathbf{h} | \mathbf{y}_i, \mathbf{P}_i^{-1}) \Rightarrow \mathbf{y}_i = \frac{1}{R} \sum_{r=1}^R \mathbf{m}_{ij}, \mathbf{P}_i^{-1} = \mathbf{L}^{-1} + \frac{1}{J} \cdot \sum_{j=1}^J (\mathbf{m}_{ij} - \mathbf{y}_i)(\mathbf{m}_{ij} - \mathbf{y}_i)^T$$

Minimum divergence estimate of speaker prior (2/2)

- NIST SRE'12 (Core task, CC2): one to over a hundred training segments per speaker, probably with content overlap among different segments for the same speaker.
 - MFCC 57, UBM 512, i-vector 400
 - By-the-book approach does not perform better than the other two approaches.
 - Results on SRE'12 show a clear benefit of MinDiv

Table2 Comparison of three speaker adaptation approaches on CC2 of NIST SRE'12 core task.

	EER (%)	minDCF10	minDCF12	
<i>By-the-book</i>	6.8953	0.6015	0.5394	Male
<i>Mean</i>	3.9395	0.4765	0.4065	
<i>MinDiv</i>	3.5746	0.4238	0.3624	
<i>By-the-book</i>	6.4646	0.6338	0.5621	Female
<i>Mean</i>	3.2145	0.5382	0.4440	
<i>MinDiv</i>	3.0597	0.5235	0.4292	

Interesting Research Topics

RAPID I-VECTOR EXTRACTION

Subspace *orthonormalizing* prior

- Consider the following informative prior

$$\mathbf{x} \sim \mathcal{N}(0, \Sigma_p) \text{ with } \Sigma_p = (\mathbf{T}^T \mathbf{T})^{-1}$$

- l-vector extraction

$$\phi = \underbrace{\mathbf{L}^{-1}}_{\substack{\uparrow \\ \left[\mathbf{T}^T \mathbf{T} + \mathbf{T}^T \mathbf{N} \mathbf{T} \right]^{-1}}} \cdot \mathbf{T}^T \tilde{\mathbf{F}} \quad \mathbf{L}^{-1} = \left[\mathbf{T}^T \mathbf{T} + \mathbf{T}^T \mathbf{N} \mathbf{T} \right]^{-1}$$

$$\begin{aligned} \phi &= \left[\mathbf{T}^T \mathbf{T} + \mathbf{T}^T \mathbf{N} \mathbf{T} \right]^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}} \\ &= \left((\mathbf{T}^T \mathbf{T}) \left[\mathbf{I} + (\mathbf{T}^T \mathbf{T})^{-1} \mathbf{T}^T \mathbf{N} \mathbf{T} \right] \right)^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}} \\ &= \left[\mathbf{I} + (\mathbf{T}^T \mathbf{T})^{-1} \mathbf{T}^T \mathbf{N} \mathbf{T} \right]^{-1} (\mathbf{T}^T \mathbf{T})^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}} \end{aligned}$$

Solving the matrix inversion (cont'd)

- Let $\mathbf{A} = (\mathbf{I} + \mathbf{N})$

$$\left[\mathbf{I} + \mathbf{N} \cdot \mathbf{T} (\mathbf{T}^T \mathbf{T})^{-1} \mathbf{T}^T \right]^{-1} = (\mathbf{I} + \mathbf{N} - \mathbf{N} \mathbf{U}_2 \mathbf{U}_2^T)^{-1} = (\mathbf{A} - \mathbf{N} \mathbf{U}_2 \mathbf{U}_2^T)^{-1}$$

- Using the matrix inversion lemma

$$(\mathbf{A} - \mathbf{N} \mathbf{U}_2 \mathbf{U}_2^T)^{-1} = \mathbf{A}^{-1} + \mathbf{A}^{-1} \mathbf{N} (\mathbf{I} - \mathbf{U}_2 \mathbf{U}_2^T \mathbf{A}^{-1} \mathbf{N})^{-1} \mathbf{U}_2 \mathbf{U}_2^T \mathbf{A}^{-1}$$

- Using again the matrix inversion identity $(\mathbf{I} + \mathbf{P} \mathbf{Q})^{-1} \mathbf{P} = \mathbf{P} (\mathbf{I} + \mathbf{Q} \mathbf{P})^{-1}$

$$\left[\mathbf{I} + \mathbf{N} \cdot \mathbf{T} (\mathbf{T}^T \mathbf{T})^{-1} \mathbf{T}^T \right]^{-1} = \mathbf{A}^{-1} + \mathbf{A}^{-1} \mathbf{N} \mathbf{U}_2 \mathbf{U}_2^T (\mathbf{I} - \mathbf{A}^{-1} \mathbf{N} \mathbf{U}_2 \mathbf{U}_2^T)^{-1} \mathbf{A}^{-1}$$

Rapid computation of i-vector

- Uniform occupancy assumption:

$$\mathbf{A}^{-1}\mathbf{N} = (\mathbf{I} + \mathbf{N})^{-1}\mathbf{N} \approx \alpha\mathbf{I} \quad \text{for } 0 \leq \alpha < 1$$

- Or equivalently

$$\frac{N_c}{1 + N_c} \approx \alpha \quad \forall c$$

- The matrix inversion can be simplified as

$$\left[\mathbf{I} + \mathbf{N} \cdot \mathbf{T}(\mathbf{T}^T\mathbf{T})^{-1}\mathbf{T}^T \right]^{-1} \approx \mathbf{A}^{-1} + \alpha \cdot \mathbf{U}_2\mathbf{U}_2^T (\mathbf{I} - \mathbf{A}^{-1}\mathbf{N}\mathbf{U}_2\mathbf{U}_2^T)^{-1} \mathbf{A}^{-1}$$

- Since $\mathbf{T} \perp \mathbf{U}_2$, the second term diminishes

$$\phi = (\mathbf{T}^T\mathbf{T})^{-1}\mathbf{T}^T \left[\mathbf{I} + \mathbf{N}\mathbf{T}(\mathbf{T}^T\mathbf{T})^{-1}\mathbf{T}^T \right]^{-1} \tilde{\mathbf{F}} \approx (\mathbf{T}^T\mathbf{T})^{-1}\mathbf{T}^T (\mathbf{I} + \mathbf{N})^{-1} \tilde{\mathbf{F}}$$

Computational complexity and memory cost

	Complexity	Memory cost	Time ratio
Baseline (slow)	$O(CFM^2 + M^3)$	$O(CFM)$	106.44
Baseline (fast)	$O(CFM + CM^2 + M^3)$	$O(CFM + CM^2)$	11.99
Proposed (exact)	$O(CFM + CM^2 + M^3)$	$O(CFM + CM^2)$	12.65
Proposed (fast)	$O(CFM)$	$O(CFM)$	1

Baseline (slow) $\phi = \left(\mathbf{I} + \sum_c N_c \mathbf{T}_c^T \mathbf{T}_c \right)^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}}$

Baseline (fast) $\phi = \left(\mathbf{I} + \sum_c N_c \mathbf{A}_c \right)^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}}$

Proposed (exact) $\phi = \left(\sum_c (N_c + 1) \mathbf{T}_c^T \mathbf{T}_c \right)^{-1} \cdot \mathbf{T}^T \tilde{\mathbf{F}}$

Proposed (fast) $\phi = \left(\mathbf{T}^T \mathbf{T} \right)^{-1} \mathbf{T}^T (\mathbf{I} + \mathbf{N})^{-1} \tilde{\mathbf{F}}$

SRE'10 core-extended (female)

- For the **tel-tel CC 5**, the relative degradation is **10.04%** in EER and **4.54%** in min DCF.

	CC1	CC2	CC3	CC4	CC5	EER	MinDCF10	CC8	CC9
baseline	2.1986	3.8313	4.3591	2.8421	3.3182	5.3826	6.3412	2.6505	2.1288
	0.3657	0.6088	0.6553	0.4593	0.5179	0.8506	0.7698	0.5199	0.2393
proposed (exact)	2.3014	3.8716	4.3417	2.7644	3.2932	5.3464	6.3280	2.6176	1.9896
	0.3680	0.6091	0.6468	0.4581	0.5192	0.8550	0.7565	0.5259	0.2212
proposed (fast)	2.4699	4.4485	4.1780	2.8356	3.6514	5.9624	7.2228	2.9741	1.8361
	0.3681	0.6507	0.7142	0.4963	0.5414	0.8349	0.7331	0.5786	0.2881
$(\mathbf{I} + \mathbf{N})^{-1} \tilde{\mathbf{F}}$	2.4144	4.5298	4.9959	3.2292	3.8124	6.2883	6.8657	2.8703	1.7414
	0.4247	0.7267	0.7856	0.5281	0.6106	0.9183	0.8013	0.5915	0.2994
$(\mathbf{I} + \mathbf{N})^{-1} \mathbf{F}$	3.5880	6.7725	6.2079	4.4367	4.6302	7.3215	9.7463	3.8837	2.2043
	0.5502	0.8591	0.8916	0.6586	0.6537	0.9457	0.8589	0.6860	0.3936

L. Xu, K. A. Lee, H. Li, and Z. Yang, “,” in *Proc. Odyssey 2016: The Speaker and Language Recognition Workshop*, 2016, pp. 47 – 52.


Kong Aik LEE

*Scientist, Institute for Infocomm Research, A*STAR, Singapore*

THANK YOU

Learning Diary

- A. Sufficient statistics are pre-whitened prior to i-vector extraction. This pre-whitening step does not change the i-vector as similar transformation is absorbed by the T matrix. Proof this.
- B. A probabilistic LDA (PLDA) model is an extension to the classical factor analysis model by tying of observed variables



Explain the rationale of tying multiple observed variables (in the context of i-vector PLDA speaker recognition system) to a single latent variable.